

Automated Transcription of a Lyric's Melody

David Branner

Hacker School, New York
20141023

Automated* Transcription of a Lyric's Melody

David Branner

Hacker School, New York
20141023

Automated* Transcription of a Lyric's Melody

(*if you've done a lot of preliminary manual work)

David Branner

Hacker School, New York
20141023

In September I reported on an experiment to transcribe the solo singing voice with automated tools.

In September I **reported** on an experiment to transcribe the solo singing voice with automated tools.

I have recordings that I want to study — specifically, I want to study how the melody reflects the organization of the words that the music is set to.

In September I **reported** on an experiment to transcribe the solo singing voice with automated tools.

I have recordings that I want to study — specifically, I want to study how the melody reflects the organization of the words that the music is set to.

I want the text, annotated with its melody. A normal musical transcription supplies the opposite of that: a melody annotated with words.

My earlier transcription (audio => MIDI => Python editing) went very badly because of tremolo in the singer's voice and reverberation in the recording.

My earlier transcription (audio => MIDI => Python editing) went very badly because of tremolo in the singer's voice and reverberation in the recording.

In the end I transcribed the piece by hand, using an open-source program called **MuseScore**.

My earlier transcription (audio => MIDI => Python editing) went very badly because of tremolo in the singer's voice and reverberation in the recording.

In the end I transcribed the piece by hand, using an open-source program called **MuseScore**. Example:

17

Pno. 

chui put - sêng - hoan chhám_ chiong - piát_
醉 不 成 歡 慘 將 別

21

Pno. 

piát - sî bông - bông_____ kang - chhim - goát_
別 時 茫 茫 江 浸 月

Initially I was pleased because MuseScore can play back the melody I transcribe, speeding the work.

Initially I was pleased because MuseScore can play back the melody I transcribe, speeding the work.

But I was in luck because it can also export the score to an open-source format called MusicXML.

Initially I was pleased because MuseScore can play back the melody I transcribe, speeding the work.

But I was in luck because it can also export the score to an open-source format called MusicXML.

This turned out to allow me to do exactly what I had originally wanted – to study the melody of each syllable, quantitatively.

Here is a single note of the 1100-odd note performance
I transcribed (the text is Taiwanese in Church *romaji*):

Here is a single note of the 1100-odd note performance
I transcribed (the text is Taiwanese in Church *romaji*):



chui

醉

Here is a single note of the 1100-odd note performance I transcribed (the text is Taiwanese in Church *romaji*):



chui
醉

and here is the XML for
that note:

Here is a single note of the 1100-odd note performance I transcribed (the text is Taiwanese in Church *romaji*):



chui
醉

and here is the XML for that note:

```
<measure number="17">
  <attributes>
    <time>
      <beats>2</beats>
      <beat-type>4</beat-type>
    </time>
  </attributes>
  <note>
    <pitch>
      <step>D</step>
      <octave>4</octave>
    </pitch>
    <duration>12</duration>
    <voice>1</voice>
```

```
    <type>quarter</type>
    <dot/>
    <stem>down</stem>
    <lyric number="1">
      <syllabic>single</syllabic>
      <text font-family="Times New
Roman">chui</text>
    </lyric>
    <lyric number="2">
      <syllabic>single</syllabic>
      <text font-family="Times New
Roman">醉</text>
    </lyric>
  </note>
```


Here is a single note of the 1100-odd note performance I transcribed (the text is Taiwanese in Church *romaji*):



chui
醉

and here is the XML for that note:

```
<measure number="17">
  <attributes>
    <time>
      <beats>2</beats>
      <beat-type>4</beat-type>
    </time>
  </attributes>
  <note>
    <pitch>
      <step>D</step>
      <octave>4</octave>
    </pitch>
    <duration>12</duration>
    <voice>1</voice>
```

```
    <type>quarter</type>
    <dot/>
    <stem>down</stem>
    <lyric number="1">
      <syllabic>single</syllabic>
      <text font-family="Times New
Roman">chui</text>
    </lyric>
    <lyric number="2">
      <syllabic>single</syllabic>
      <text font-family="Times New
Roman">醉</text>
    </lyric>
  </note>
```

That's exactly what I need to turn

**That's exactly what I need to turn
words set to music (traditional score-notation)**

That's exactly what I need to turn
words set to music (traditional score-notation)
into


That's exactly what I need to turn
words set to music (traditional score-notation)
into
melody on each syllable.

That's exactly what I need to turn
words set to music (traditional score-notation)
into
melody on each syllable.

The point is to have a representation of the melody
organized by words, rather than the other way around.


Another example:

Another example:



bông
茫

Another example:




bông
茫

```
<note>
  <pitch>
    <step>E</step>
    <octave>4</octave>
  </pitch>
  <duration>6</duration>
  <voice>1</voice>
  <type>eighth</type>
  <dot/>
  <stem>down</stem>
  <beam number="1">begin</beam>
  <notations>
    <slur type="start"
number="1"/>
  </notations>
  <lyric number="1">
    <syllabic>end</syllabic>
    <text font-family="Times
New Roman">bông</text>
  </lyric>
  <lyric number="2">
    <syllabic>single</syllabic>
```

```
    <text>茫</text>
  </lyric>
</note>
<note>
  <pitch>
    <step>G</step>
    <octave>4</octave>
  </pitch>
  <duration>2</duration>
  <voice>1</voice>
  <type>16th</type>
  <stem>down</stem>
  <beam number="1">continue</
beam>
  <beam number="2">forward
hook</beam>
</note>
<note>
  <pitch>
    <step>E</step>
    <octave>4</octave>
  </pitch>
```

```
  <duration>4</duration>
  <voice>1</voice>
  <type>eighth</type>
  <stem>down</stem>
  <beam number="1">continue</
beam>
</note>
<note>
  <pitch>
    <step>G</step>
    <octave>4</octave>
  </pitch>
  <duration>4</duration>
  <voice>1</voice>
  <type>eighth</type>
  <stem>down</stem>
  <beam number="1">end</beam>
  <notations>
    <slur type="stop"
number="1"/>
  </notations>
</note>
```

Another example:



bông
茫

```
<note>
  <pitch>
    <step>E</step>
    <octave>4</octave>
  </pitch>
</duration>6</duration>
<voice>1</voice>
<type>eighth</type>
<dot/>
<stem>down</stem>
<beam number="1">begin</beam>
<notations>
  <slur type="start"
number="1"/>
</notations>
<lyric number="1">
  <syllabic>end</syllabic>
  <text font-family="Times
New Roman">bông</text>
</lyric>
<lyric number="2">
  <syllabic>single</syllabic>
```

```
<text>茫</text>
</lyric>
</note>
<note>
  <pitch>
    <step>G</step>
    <octave>4</octave>
  </pitch>
</duration>2</duration>
<voice>1</voice>
<type>16th</type>
<stem>down</stem>
<beam number="1">continue</
beam>
  <beam number="2">forward
hook</beam>
</note>
<note>
  <pitch>
    <step>E</step>
    <octave>4</octave>
  </pitch>
```

```
</duration>4</duration>
<voice>1</voice>
<type>eighth</type>
<stem>down</stem>
<beam number="1">continue</
beam>
</note>
<note>
  <pitch>
    <step>G</step>
    <octave>4</octave>
  </pitch>
</duration>4</duration>
<voice>1</voice>
<type>eighth</type>
<stem>down</stem>
<beam number="1">end</beam>
<notations>
  <slur type="stop"
number="1"/>
</notations>
</note>
```

It is easy to parse the XML using the `lxml` library; `lxml`'s `root.xpath` method allows me request all the `note` elements as a generator, and I can step through the generator's output, identifying syllables and then retrieving the melody to assigned to those syllables.

Important edge cases:

Important edge cases:

1. rests — they have no pitch content; also, since they don't correspond to syllables, consecutive rests can be collapsed into one;

Important edge cases:

1. rests — they have no pitch content; also, since they don't correspond to syllables, consecutive rests can be collapsed into one;
2. notes tied (but not slurred) together — since they correspond to one or part of one syllable and are all the same pitch, they can be collapsed into one;

Important edge cases:

1. rests — they have no pitch content; also, since they don't correspond to syllables, consecutive rests can be collapsed into one;
2. notes tied (but not slurred) together — since they correspond to one or part of one syllable and are all the same pitch, they can be collapsed into one;
3. melisma — more than one note per vowel or syllable — represented as a series of notes on one syllable.

Output of the first note above:

Output of the first note above:



chuí
醉

Output of the first note above:



chuí
醉

```
...  
('chuí',  
  [{'pitch_data': {'octave': '4', 'step': 'D'},  
    'lyric_2': {'syllabic': 'single', 'text': '醉'},  
    'duration': 12}]),  
...
```

Output of the more melismatic second example above:

Output of the more melismatic second example above:



bông _____
茫

Output of the more melismatic second example above:



bông
茫

```
...  
('bông',  
  [{'pitch_data': {'octave': '4', 'step': 'E'}},  
    'lyric_2': {'syllabic': 'single', 'text': '茫'},  
    'duration': 6},  
  {'pitch_data': {'octave': '4', 'step': 'G'}}, 'duration': 2},  
  {'pitch_data': {'octave': '4', 'step': 'E'}}, 'duration': 4},  
  {'pitch_data': {'octave': '4', 'step': 'G'}}, 'duration': 4}]),  
...)
```

Output of the more melismatic second example above:



bông
茫

```
...  
('bông',  
  [{'pitch_data': {'octave': '4', 'step': 'E'}},  
   'lyric_2': {'syllabic': 'single', 'text': '茫'},  
   'duration': 6},  
  {'pitch_data': {'octave': '4', 'step': 'G'}}, 'duration': 2},  
  {'pitch_data': {'octave': '4', 'step': 'E'}}, 'duration': 4},  
  {'pitch_data': {'octave': '4', 'step': 'G'}}, 'duration': 4}]),  
...  
...
```

Final output is list of tuples: [(syllable, [note-dicts]), ...].

Final output is list of tuples: [(syllable, [note-dicts]), ...].

Next steps:

Final output is list of tuples: [(syllable, [note-dicts]), ...].

Next steps:

1. Conversion of 'octave' and 'step' attributes of 'pitch_data' to *MIDI* pitch numbers.

Final output is list of tuples: [(syllable, [note-dicts]), ...].

Next steps:

1. Conversion of 'octave' and 'step' attributes of 'pitch_data' to *MIDI* pitch numbers.
2. Analysis of the melodic values of syllables and language phrases.

End